

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

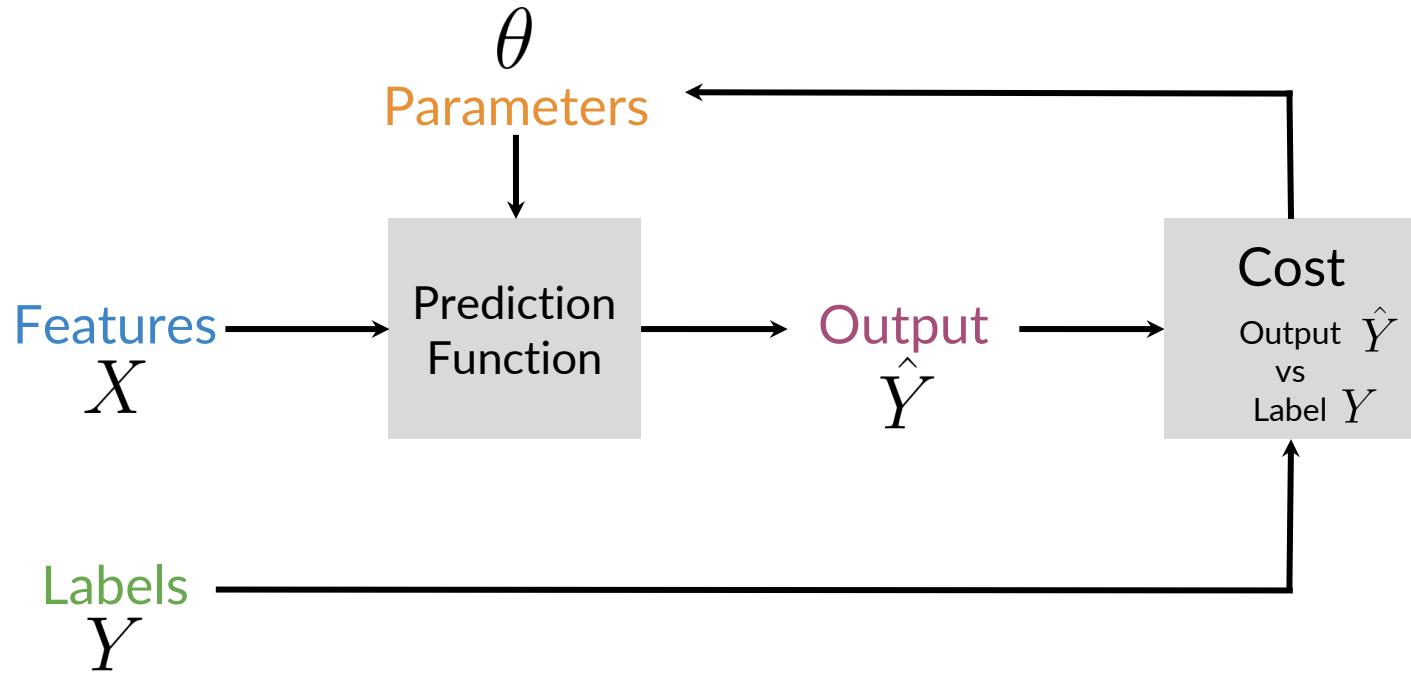
# Supervised ML and Sentiment Analysis

---

# Outline

- Review Supervised ML
- Build your own tweet classifier!

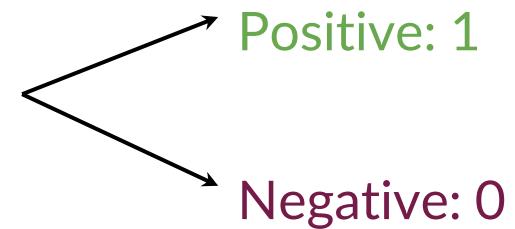
# Supervised ML (training)



# Sentiment analysis

Tweet I am happy because I am learning NLP

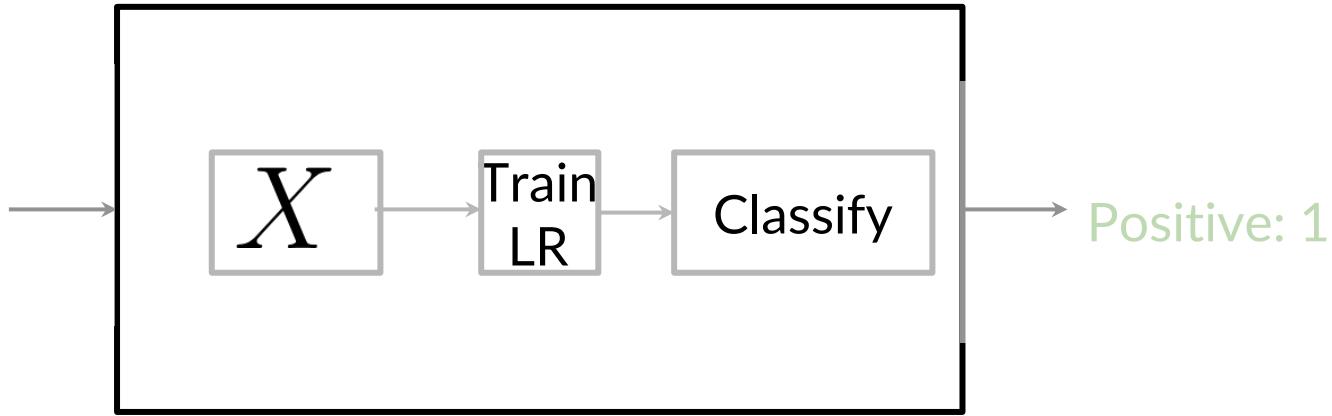
:



Logistic regression

# Sentiment analysis

I am happy  
because I am  
learning NLP



# Summary

- Features, Labels → Train → Predict
- Extract features → Train+LR → Predict sentiment



deeplearning.ai

# Vocabulary and Feature Extraction

---

# Outline

- Vocabulary
- Feature extraction
- Sparse representations and some of their issues

# Vocabulary

Tweets:

[tweet\_1, tweet\_2, ..., tweet\_m]



I am happy because I am learning

NLP

...

...

I hated the movie

$V =$

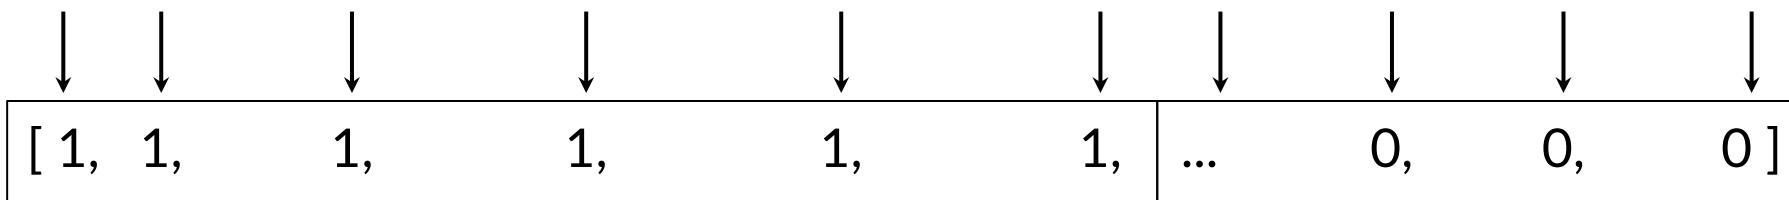
[ I, am, happy, because, learning, NLP, ...      hated, the, movie ]

*all unique words  
in all tweets*

# Feature extraction

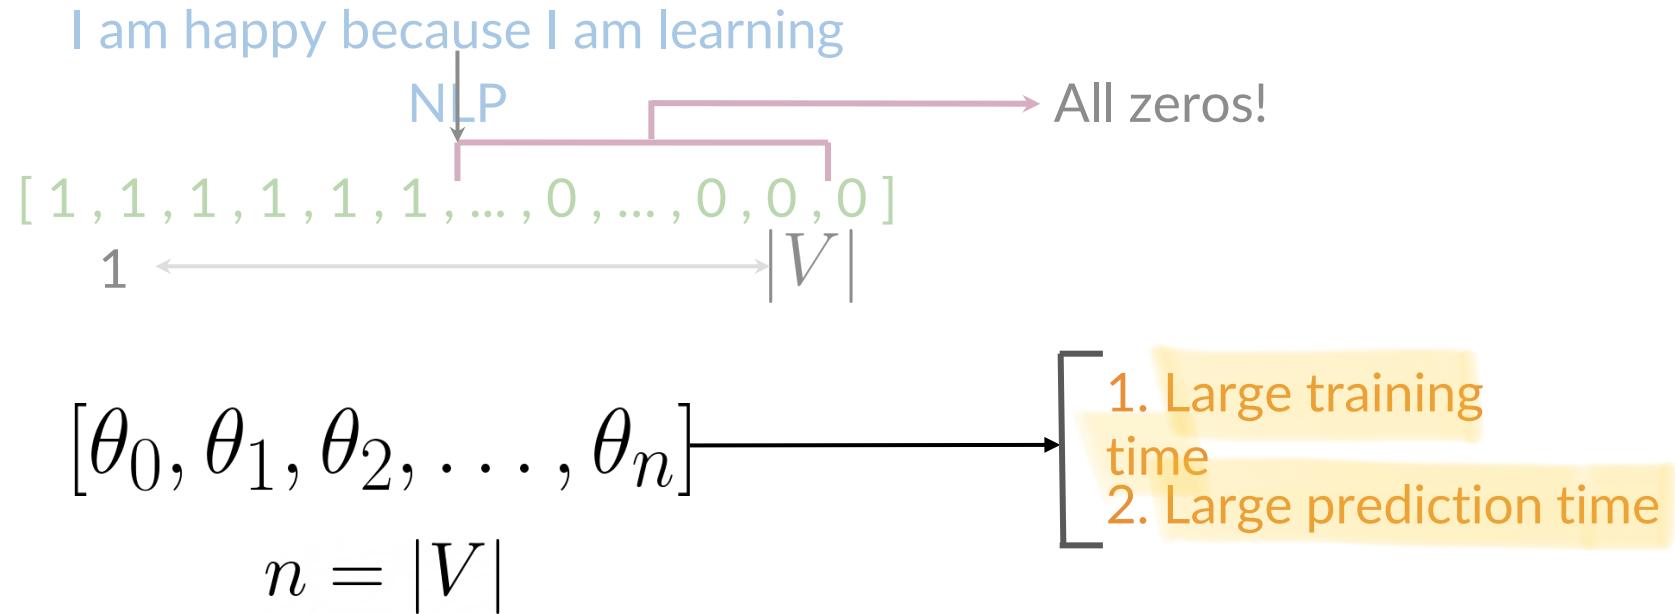
I am happy because I am learning NLP

[ I , am, happy, because, learning, NLP, ... hated, the, movie ]



A lot of zeros! That's a sparse representation.

# Problems with sparse representations



# Summary

- Vocabulary: set of unique words
- Vocabulary, Text → [1 ..... 0 ..... 1 .. 0 .. 1 .. 0]
- Sparse representations are problematic for training and prediction times



deeplearning.ai

# Negative and Positive Frequencies

# Outline

- Populate your vocabulary with a frequency count for each class

# Positive and negative counts

## Corpus

I am happy because I am learning NLP

I am happy

I am sad, I am not learning NLP

I am sad

## Vocabulary

I

am

happy

because

learning

NLP

sad

not

# Positive and negative counts

Positive tweets

I am happy because I am learning  
NLP  
I am happy

Negative tweets

I am sad, I am not learning NLP  
I am sad

# Positive and negative counts

## Positive tweets

I am happy because I am learning  
NLP  
I am happy

Vocabulary    PosFreq (1)

---

I

am

happy

2

because

learning

NLP

sad

not

0

---

# Positive and negative counts

Vocabulary	NegFreq (0)
I	
am	3
happy	
because	
learning	
NLP	
sad	
not	1

Negative tweets

I am sad, I am not learning NLP  
I am sad

freq dict

## Word frequency in classes

Vocabulary	PosFreq (1)	NegFreq (0)
I	3	3
am	3	3
happy	2	0
because	1	0
learning	1	1
NLP	1	1
sad	0	2
not	0	1

*freqs*: dictionary mapping from (word, class) to frequency

# Summary

- Divide tweet corpus into two classes: positive and negative
  - Count each time each word appears in either class
- Feature extraction for training and prediction!



deeplearning.ai

# Feature extraction with frequencies

# Outline

- Extract features from your frequencies dictionary to create a features vector

# Word frequency in classes

Vocabulary	PosFreq (1)	NegFreq (0)
I	3	3
am	3	3
happy	2	0
because	1	0
learning	1	1
NLP	1	1
sad	0	2
not	0	1

*freqs*: dictionary mapping from (word, class) to frequency

# Feature extraction

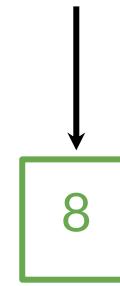
**freqs**: dictionary mapping from (word, class) to frequency

# Feature extraction

Vocabulary	PosFreq (1)
I	<u>3</u>
am	<u>3</u>
happy	2
because	1
learning	<u>1</u>
NLP	<u>1</u>
sad	<u>0</u>
not	<u>0</u>

I am sad, I am not learning NLP

$$X_m = [1, \sum_w freqs(w, 1), \sum_w freqs(w, 0)]$$

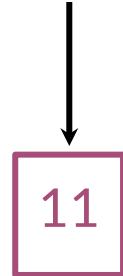


# Feature extraction

Vocabulary	NegFreq (0)
I	<u>3</u>
am	<u>3</u>
happy	0
because	0
learning	<u>1</u>
NLP	<u>1</u>
sad	<u>2</u>
not	<u>1</u>

I am sad, I am not learning NLP

$$X_m = [1, \sum_w freqs(w, 1), \sum_w freqs(w, 0)]$$



# Feature extraction

I am sad, I am not learning NLP

$$X_m = [1, \sum_w freqs(w, 1), \sum_w freqs(w, 0)]$$

↓

Now we have 3  
features instead of  $\sqrt{|\text{vocab}|}$   
 $X_m = [1, 8, 11]$  a vector  
of tweet  
features.

# Summary

- Dictionary mapping (word,class) to frequencies

$$X_m = [1, \sum_w freqs(w, 1), \sum_w freqs(w, 0)]$$

- Cleaning unimportant information from your tweets



deeplearning.ai

# Preprocessing

---

# Outline

- Removing stopwords, punctuation, handles and URLs
- Stemming
- Lowercasing

# Preprocessing: stop words and punctuation

@YMourri and @AndrewYNg are tuning a GREAT AI model at <https://deeplearning.ai!!!>

Stop words	Punctuation
and	,
is	.
are	:
at	!
has	"
for	'
a	

# Preprocessing: stop words and punctuation

@YMourri ~~and~~ @AndrewYNg ~~are~~  
tuning ~~a~~ GREAT AI model ~~at~~  
[https://deeplearning.ai!!!](https://deeplearning.ai)

@YMourri @AndrewYNg tuning  
GREAT AI model  
[https://deeplearning.ai!!!](https://deeplearning.ai)

<u>Stop words</u>
<u>and</u>
<u>is</u>
<u>are</u>
<u>at</u>
<u>has</u>
<u>for</u>
<u>a</u>

<u>Punctuation</u>
,
.
:
!
"
'

# Preprocessing: stop words and punctuation

@YMourri @AndrewYNg tuning  
GREAT AI model  
<https://deeplearning.ai!!!>

@YMourri @AndrewYNg tuning  
GREAT AI model  
<https://deeplearning.ai>

Stop words	Punctuation
and	,
is	.
a	:
at	!
has	"
for	'
of	

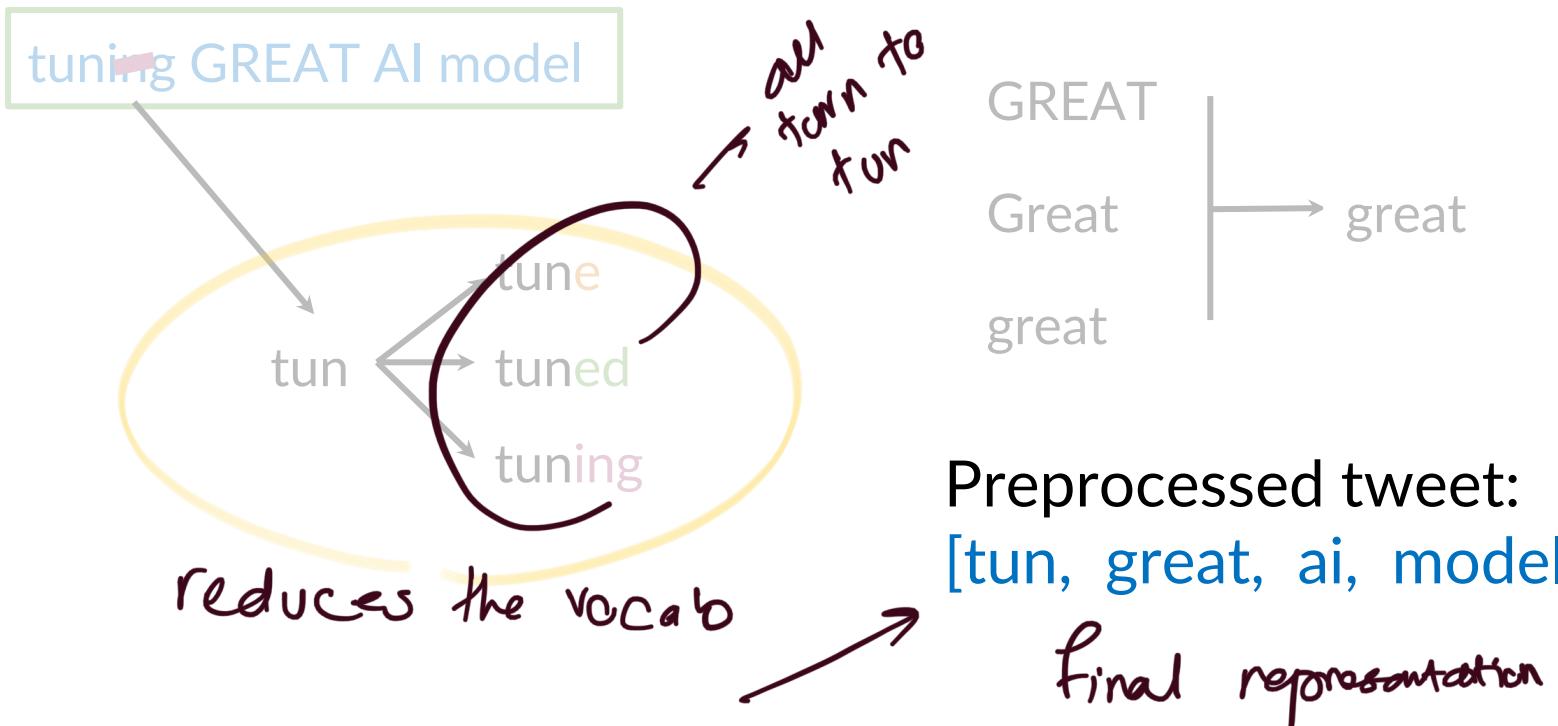
# Preprocessing: Handles and URLs

~~@YMurri @AndrewYNg tuning GREAT AI  
model~~

<https://deeplearning.ai>

tuning GREAT AI model

# Preprocessing: Stemming and lowercasing



# Summary

- Stop words, punctuation, handles and URLs
- Stemming
- Lowercasing
- Less unnecessary info → Better times



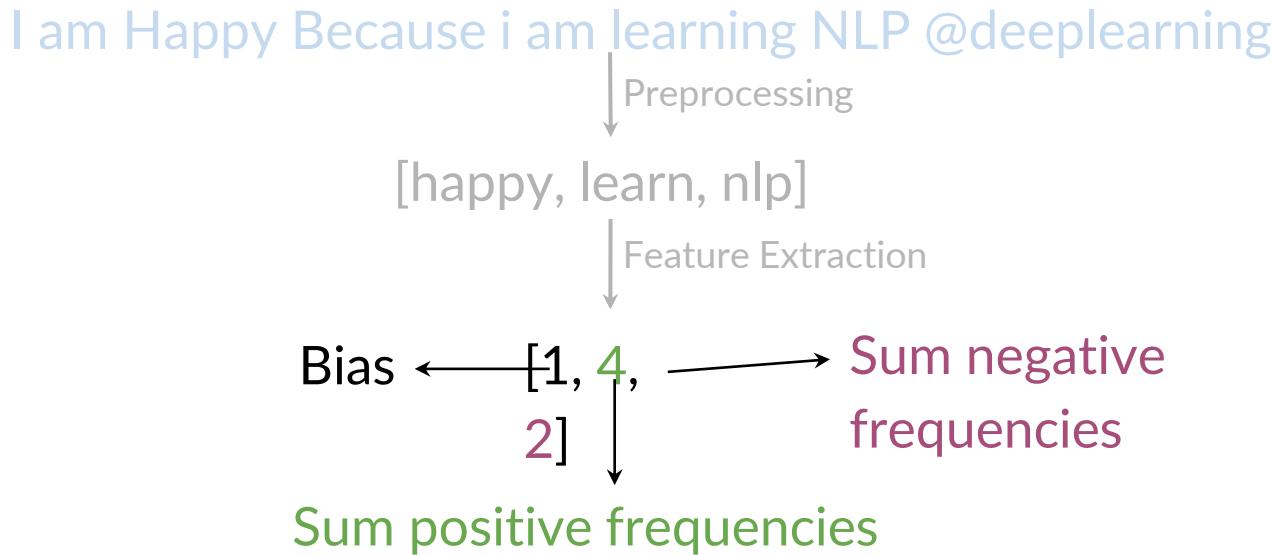
deeplearning.ai

Putting it all  
together

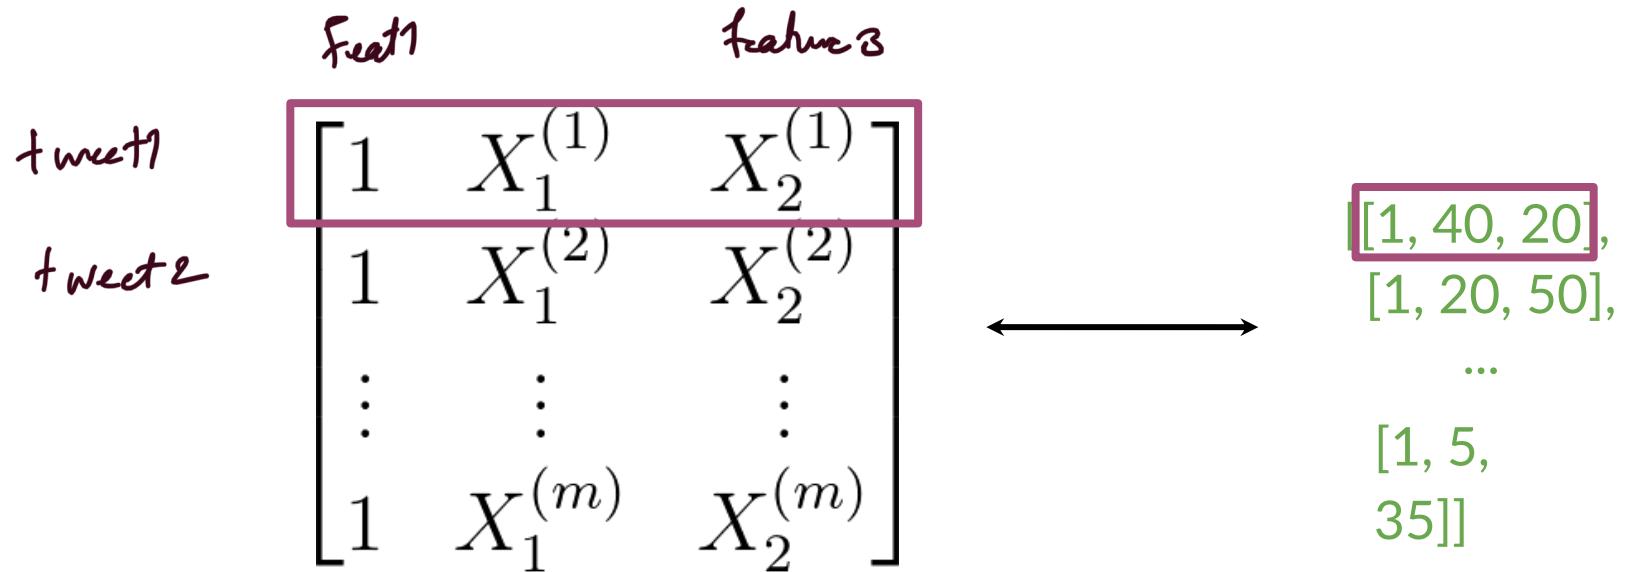
# Outline

- Generalize the process
- How to code it!

# General overview



# General overview



# General Implementation

```
freqs = build_freqs(tweets,labels) #Build frequencies dictionary  
X = np.zeros((m,3)) #Initialize matrix X  
  
for i in range(m): #For every tweet  
    p_tweet = process_tweet(tweets[i]) #Process tweet  
    X[i,:] = extract_features(p_tweet,freqs) #Extract Features
```

# Summary

- Implement the feature extraction algorithm for your entire set of tweets
- Almost ready to train!



deeplearning.ai

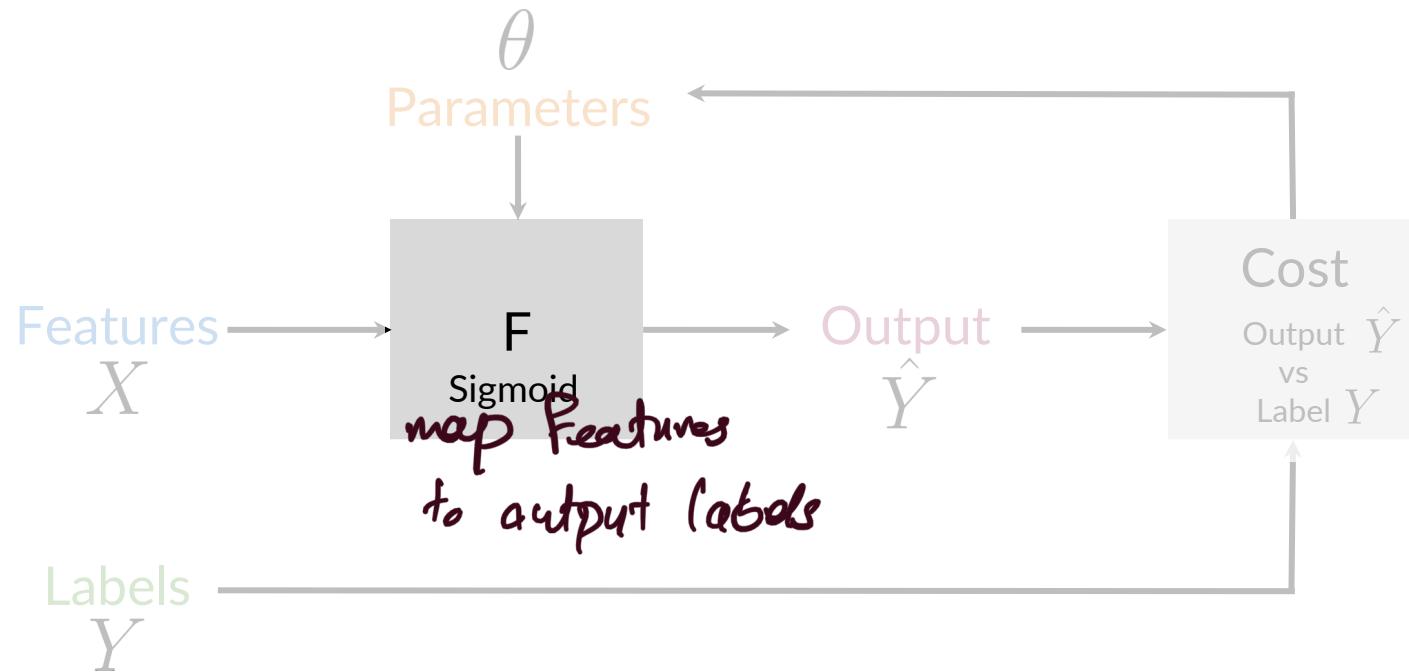
# Logistic Regression Overview

---

# Outline

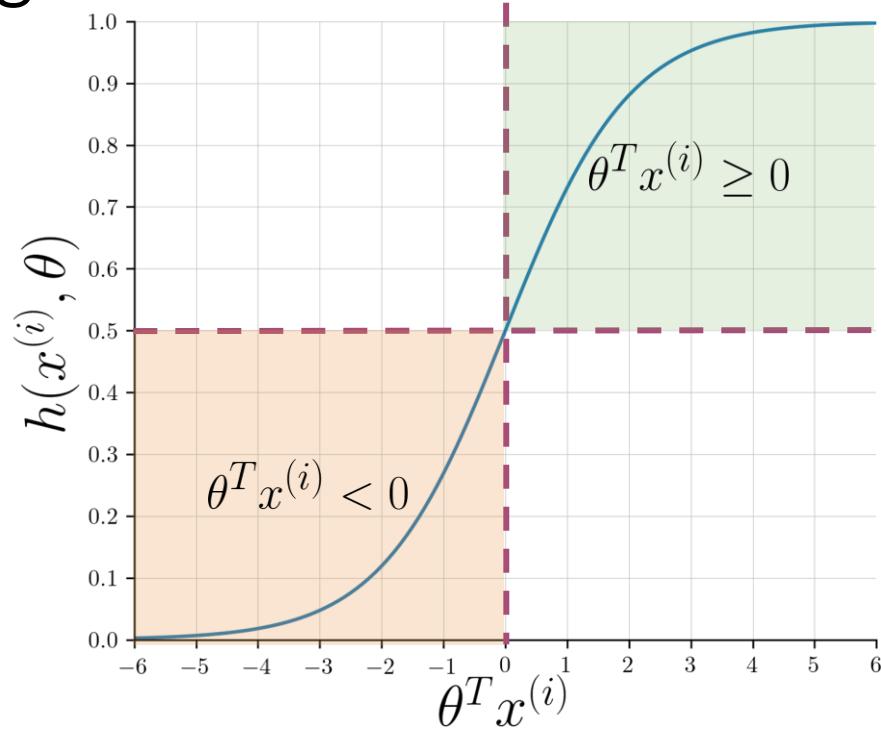
- Supervised learning and logistic regression
- Sigmoid function

# Overview of logistic regression



# Overview of logistic regression

$$h(x^{(i)}, \theta) = \frac{1}{1 + e^{-\theta^T x^{(i)}}}$$

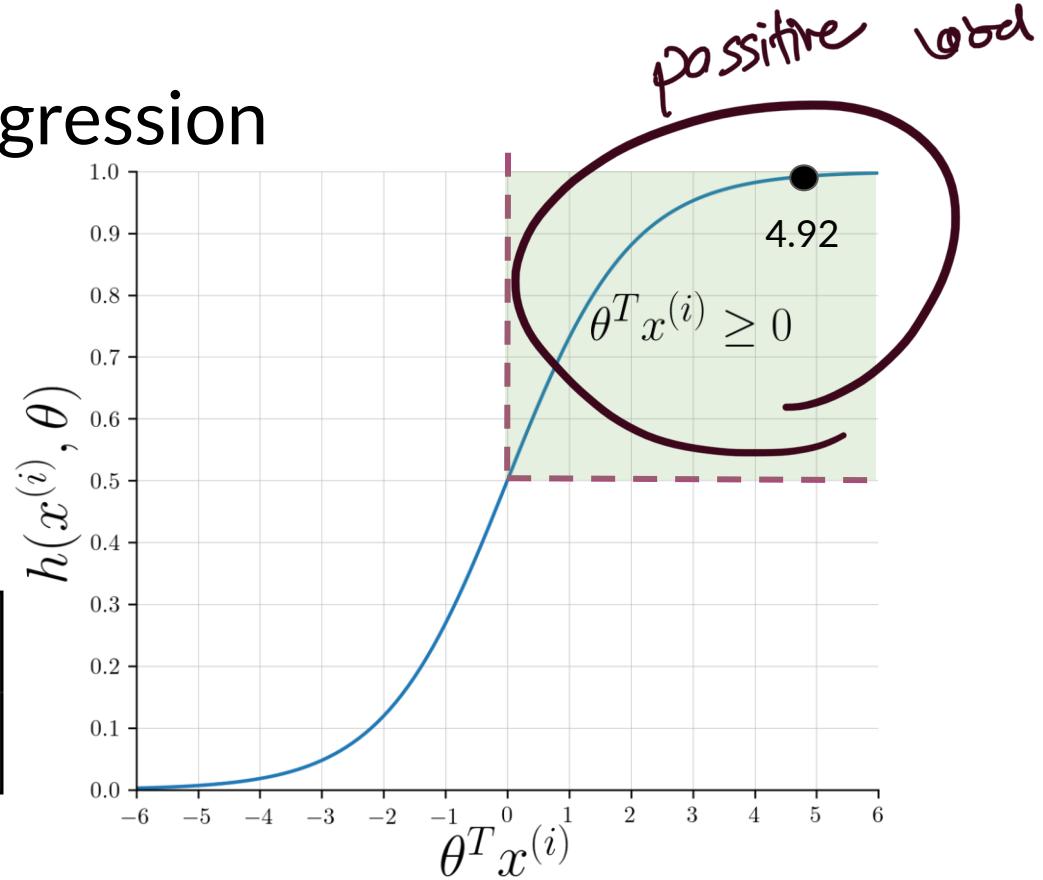


# Overview of logistic regression

@YMourri and  
@AndrewYNg are tuning a  
GREAT AI model

[tun, ai, great,  
model]

$$x^{(i)} = \begin{bmatrix} 1 \\ 3476 \\ 245 \end{bmatrix} \quad \theta = \begin{bmatrix} 0.00003 \\ 0.00150 \\ -0.00120 \end{bmatrix}$$



# Summary

- Sigmoid function
- $\theta^T x^{(i)} \geq 0 \longrightarrow h(x^{(i)}, \theta) \geq 0.5$  , positive
- $\theta^T x^{(i)} < 0 \longrightarrow h(x^{(i)}, \theta) < 0.5$  , negative



deeplearning.ai

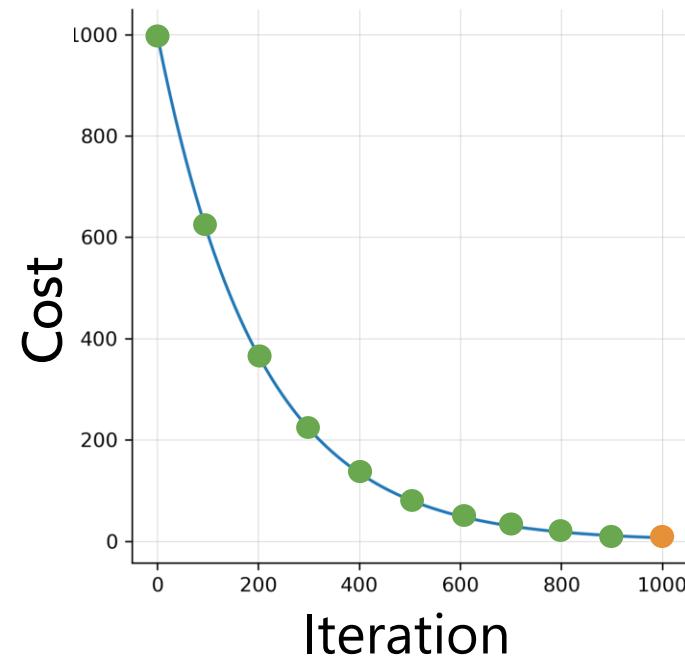
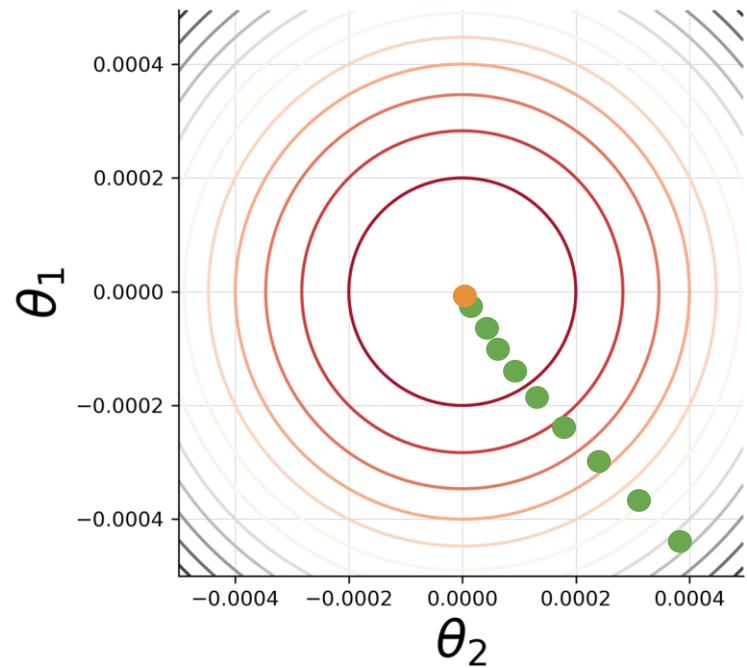
# Logistic Regression: Training

---

# Outline

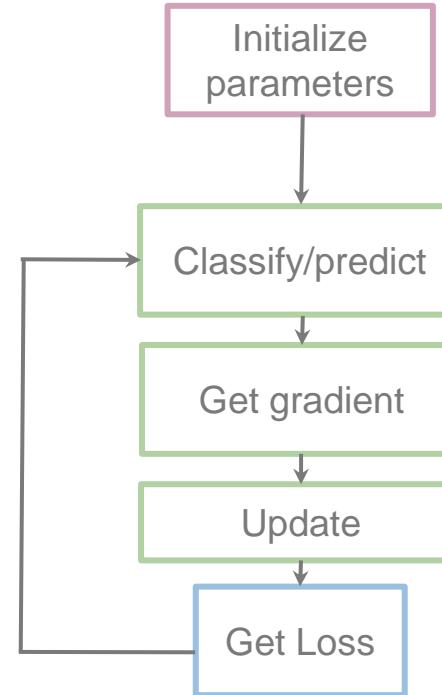
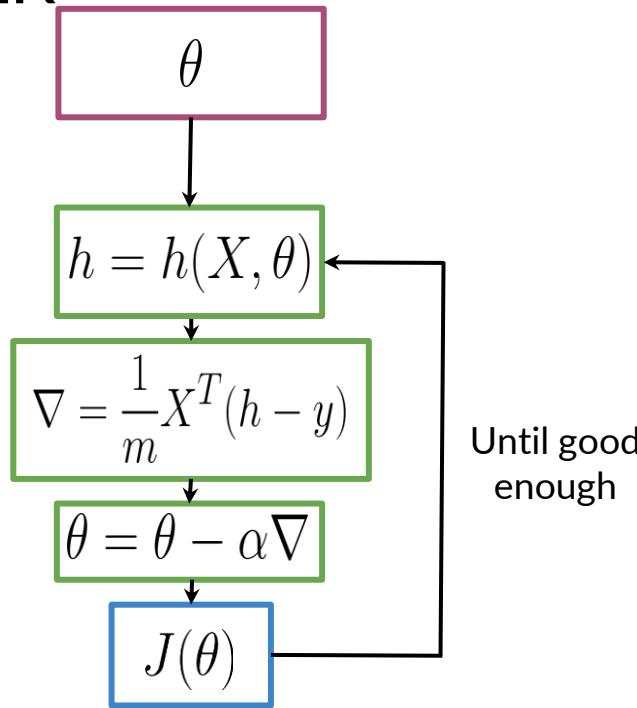
- Review the steps in the training process
- Overview of gradient descent

# Training LR



*gradient descent*

# Training LR



# Summary

- Visualize how gradient descent works
  - Use gradient descent to train your logistic regression classifier
- Compute the accuracy of your model



deeplearning.ai

# Logistic Regression: Testing

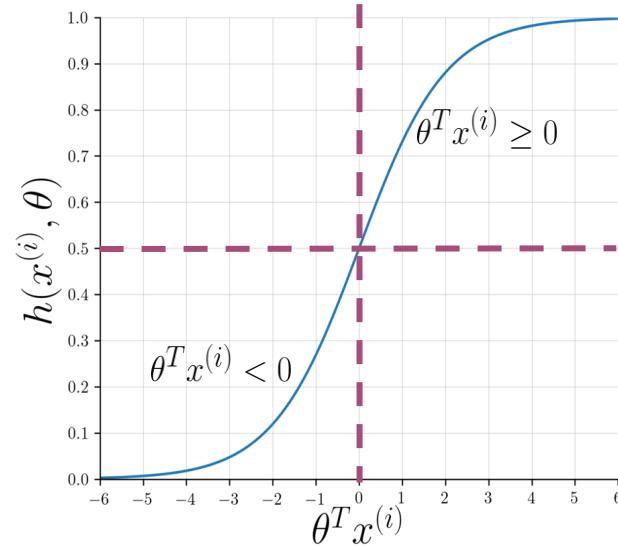
---

# Outline

- Using your validation set to compute model accuracy
- What the accuracy metric means

# Testing logistic regression

- $X_{val} \ Y_{val} \ \theta$   
 $h(X_{val}, \theta)$
- $pred = h(X_{val}, \theta) \geq 0.5$



# Testing logistic regression

- $X_{val} \ Y_{val} \ \theta$

$$h(X_{val}, \theta)$$

$$pred = h(X_{val}, \theta) \geq 0.5$$

$$\begin{bmatrix} 0.3 \\ 0.8 \\ 0.5 \\ \vdots \\ h_m \end{bmatrix} \geq 0.5 = \begin{bmatrix} \underline{0.3 \geq 0.5} \\ \underline{0.8 \geq 0.5} \\ \underline{0.5 \geq 0.5} \\ \vdots \\ pred_m \geq 0.5 \end{bmatrix} = \begin{bmatrix} \underline{0} \\ \underline{1} \\ \underline{1} \\ \vdots \\ pred_m \end{bmatrix}$$

# Testing logistic regression

- $X_{val}$   $Y_{val}$   $\theta$   
 $h(X_{val}, \theta)$

$$pred = h(X_{val}, \theta) \geq 0.5$$

$$\sum_{i=1}^m \frac{(pred^{(i)} == y_{val}^{(i)})}{m}$$

$$\begin{bmatrix} 0 \\ 1 \\ 1 \\ \vdots \\ pred_m \end{bmatrix} == \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ Y_{val_m} \end{bmatrix}$$

$$\begin{bmatrix} 1 \\ 0 \\ 1 \\ \vdots \\ pred_m == Y_{val_m} \end{bmatrix}$$

# Testing logistic regression

$$Y_{val} = \begin{bmatrix} 0 \\ 1 \\ \underline{1} \\ 0 \\ 1 \end{bmatrix} \quad pred = \begin{bmatrix} 0 \\ 1 \\ \underline{0} \\ 0 \\ 1 \end{bmatrix}$$

$$(Y_{val} == pred) = \begin{bmatrix} 1 \\ 1 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

$$\text{accuracy} = \frac{4}{5} = 0.8$$

# Summary

- $X_{val} \ Y_{val}$  → Performance on unseen data
- Accuracy →  $\sum_{i=1}^m \frac{(pred^{(i)} == y_{val}^{(i)})}{m}$

To improve model: step size, number of iterations, regularization, new features, etc.



deeplearning.ai

# Logistic Regression: Cost Function

---

# Outline

- Overview of the logistic cost function, AKA the binary cross-entropy function

# Cost function for logistic regression

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

arg over

all training samples

(# m )

# Cost function for logistic regression

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

# Cost function for logistic regression

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

label      prediction

$y^{(i)}$      $h(x^{(i)}, \theta)$

0	any	0	} ok prediction
1	0.99	~0	
1	~0	-inf	wrong prediction

# Cost function for logistic regression

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

*negative*                                   *negative*

$$\begin{array}{ccc} y^{(i)} & h(x^{(i)}, \theta) \\ \hline \end{array}$$

$$1 \quad \text{any} \quad 0$$

$$0 \quad 0.01 \quad \sim 0$$

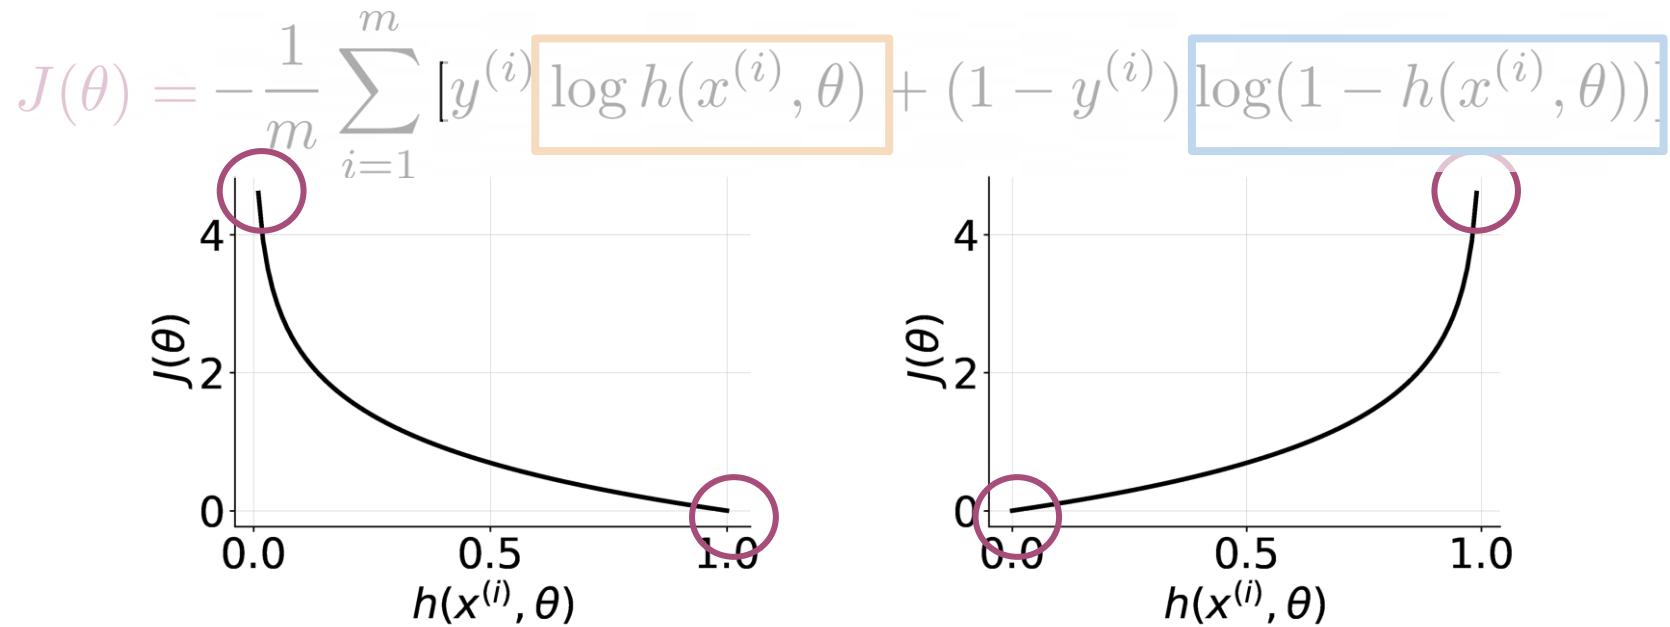
$$0 \quad \sim 1 \quad -\infty$$

→ Wrong prediction

# Cost function for logistic regression

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

# Cost function for logistic regression



# Summary

- Strong disagreement = high cost
- Strong agreement = low cost
- Aim for the lowest cost!